

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
```

```
# Load MNIST dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
# Normalize images
X_train = X_train / 255.0
X_test = X_test / 255.0
```


```
# Reshape to add channel dimension
X_train = X_train.reshape(-1, 28, 28, 1)
X_test = X_test.reshape(-1, 28, 28, 1)
```

```
# One-hot encode labels
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)
```

 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 ————— 0s 0us/step


```
# Build the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Dropout(0.3),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Dropout(0.3),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.4),
    Dense(10, activation='softmax') # 10 classes for digits (0-9)
])
```

```
# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

 /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape` to `input\_shape` in the constructor of `Conv2D` or `Conv3D` layers. It is deprecated and will be removed in a future version. Use `input\_shape` in the `compile` method instead.  
super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

```
# Train the model
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.2)
```

```
# Evaluate the model
test_loss, test_acc = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_acc:.2f}")
```

 Epoch 1/10  
1500/1500 ————— 48s 31ms/step - accuracy: 0.8345 - loss: 0.5107 - val\_accuracy: 0.9798 - val\_loss: 0.0687  
Epoch 2/10  
1500/1500 ————— 45s 30ms/step - accuracy: 0.9670 - loss: 0.1096 - val\_accuracy: 0.9858 - val\_loss: 0.0466  
Epoch 3/10  
1500/1500 ————— 82s 30ms/step - accuracy: 0.9785 - loss: 0.0746 - val\_accuracy: 0.9884 - val\_loss: 0.0404  
Epoch 4/10  
1500/1500 ————— 82s 30ms/step - accuracy: 0.9806 - loss: 0.0641 - val\_accuracy: 0.9898 - val\_loss: 0.0342  
Epoch 5/10  
1500/1500 ————— 81s 30ms/step - accuracy: 0.9819 - loss: 0.0554 - val\_accuracy: 0.9896 - val\_loss: 0.0342  
Epoch 6/10  
1500/1500 ————— 48s 32ms/step - accuracy: 0.9838 - loss: 0.0518 - val\_accuracy: 0.9908 - val\_loss: 0.0301  
Epoch 7/10  
1500/1500 ————— 81s 31ms/step - accuracy: 0.9851 - loss: 0.0484 - val\_accuracy: 0.9903 - val\_loss: 0.0336  
Epoch 8/10  
1500/1500 ————— 80s 30ms/step - accuracy: 0.9856 - loss: 0.0452 - val\_accuracy: 0.9898 - val\_loss: 0.0328  
Epoch 9/10  
1500/1500 ————— 85s 32ms/step - accuracy: 0.9870 - loss: 0.0405 - val\_accuracy: 0.9913 - val\_loss: 0.0328  
Epoch 10/10  
1500/1500 ————— 80s 30ms/step - accuracy: 0.9887 - loss: 0.0372 - val\_accuracy: 0.9930 - val\_loss: 0.0281

313/313 ————— 2s 7ms/step - accuracy: 0.9902 - loss: 0.0266  
 Test Accuracy: 0.99

```
# Save the trained model
model.save("handwritten_character_recognition.h5")
```

⚠ WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is consi

handwritten digits prediction using deep learning

```
from sklearn.datasets import fetch_openml
mnist=fetch_openml("mnist_784")

import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# Load the MNIST dataset
mnist = fetch_openml('mnist_784', version=1)
X, y = mnist.data, mnist.target

# Ensure that the data is in the correct format
X = X.astype(np.float32) # Convert pixel values to float32
y = y.astype(np.int8)    # Convert target values to integers

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Initialize the MLP classifier
mlp = MLPClassifier(hidden_layer_sizes=(64, 64),
                    max_iter=50,
                    alpha=1e-4,
                    solver='adam',
                    random_state=42,
                    verbose=10,
                    tol=1e-4)

# Train the model
mlp.fit(X_train, y_train)

# Predict on the test data
y_pred = mlp.predict(X_test)

# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")

# Visualize some predictions
fig, axes = plt.subplots(2, 5, figsize=(10, 5))
for ax, image, prediction, true_label in zip(axes.ravel(), X_test.to_numpy(), y_pred, y_test):
    ax.set_axis_off()
    image = image.reshape(28, 28)
    ax.imshow(image, cmap='gray')
    ax.set_title(f"Pred: {prediction}, True: {true_label}")

plt.tight_layout()
plt.show()
```

```

Iteration 1, loss = 2.10292369
Iteration 2, loss = 0.63468546
Iteration 3, loss = 0.41170029
Iteration 4, loss = 0.32570292
Iteration 5, loss = 0.27403755
Iteration 6, loss = 0.24147559
Iteration 7, loss = 0.21699058
Iteration 8, loss = 0.19487723
Iteration 9, loss = 0.17771913
Iteration 10, loss = 0.16357822
Iteration 11, loss = 0.15891865
Iteration 12, loss = 0.14965729
Iteration 13, loss = 0.12972488
Iteration 14, loss = 0.13554853
Iteration 15, loss = 0.13080247
Iteration 16, loss = 0.13119637
Iteration 17, loss = 0.11904129
Iteration 18, loss = 0.11818511
Iteration 19, loss = 0.11466374
Iteration 20, loss = 0.11052036
Iteration 21, loss = 0.10766985
Iteration 22, loss = 0.10104356
Iteration 23, loss = 0.09436340
Iteration 24, loss = 0.08939378
Iteration 25, loss = 0.09473639
Iteration 26, loss = 0.08972806
Iteration 27, loss = 0.08268881
Iteration 28, loss = 0.08483314
Iteration 29, loss = 0.08084884
Iteration 30, loss = 0.07425225
Iteration 31, loss = 0.06775623
Iteration 32, loss = 0.06304130
Iteration 33, loss = 0.07305904
Iteration 34, loss = 0.06126501
Iteration 35, loss = 0.05835775
Iteration 36, loss = 0.06709495
Iteration 37, loss = 0.06276928
Iteration 38, loss = 0.05704420
Iteration 39, loss = 0.05660724
Iteration 40, loss = 0.04841852
Iteration 41, loss = 0.05089388
Iteration 42, loss = 0.05344457
Iteration 43, loss = 0.05155276
Iteration 44, loss = 0.04803985
Iteration 45, loss = 0.04400076
Iteration 46, loss = 0.04093383
Iteration 47, loss = 0.03991760
Iteration 48, loss = 0.03962692
Iteration 49, loss = 0.04080373
Iteration 50, loss = 0.04110693

```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	1343
1	0.98	0.98	0.98	1600
2	0.96	0.94	0.95	1380
3	0.93	0.95	0.94	1433
4	0.94	0.97	0.96	1295
5	0.97	0.93	0.95	1273
6	0.98	0.97	0.98	1396
7	0.95	0.97	0.96	1503
8	0.94	0.93	0.94	1357
9	0.95	0.94	0.94	1420
accuracy			0.96	14000
macro avg	0.96	0.96	0.96	14000
weighted avg	0.96	0.96	0.96	14000

Confusion Matrix:

```

[[1314 1 5 2 2 1 2 3 10 3]
 [ 0 1565 3 3 1 1 3 7 12 5]
 [ 4 8 1304 12 8 3 3 24 11 3]
 [ 0 2 12 1364 2 14 0 9 12 18]
 [ 2 1 2 0 1260 0 3 10 3 14]
 [ 3 6 5 41 4 1178 12 4 15 5]
 [ 6 5 2 2 11 8 1356 1 3 2]
 [ 4 2 8 2 11 1 0 1458 2 15]
 [ 2 9 17 29 5 9 2 8 1266 10]
 [ 8 2 3 9 36 4 1 17 9 1331]]

```

Accuracy: 95.69%

/usr/local/lib/python3.11/dist-packages/sklearn/neural\_network/\_multilayer\_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: WARNING: The optimization process terminated with a warning: warnings.warn(

Pred: 8, True: 8

Pred: 4, True: 4

Pred: 5, True: 8

Pred: 7, True: 7

Pred: 7, True: 7