

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
```

```
# Load MNIST dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
# Normalize images
X_train = X_train / 255.0
X_test = X_test / 255.0
```

```
# Reshape to add channel dimension
X_train = X_train.reshape(-1, 28, 28, 1)
X_test = X_test.reshape(-1, 28, 28, 1)
```

```
# One-hot encode labels
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 ————— 0s 0us/step

```
# Build the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Dropout(0.3),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Dropout(0.3),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.4),
    Dense(10, activation='softmax') # 10 classes for digits (0-9)
])
```


```
# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` to `input_shape` in the constructor of `Conv2D` or `Conv3D` layers. It is deprecated and will be removed in a future version. Use `input_shape` in the `compile` method instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)


```
# Train the model
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.2)
```

```
# Evaluate the model
test_loss, test_acc = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_acc:.2f}")
```

Epoch 1/10
1500/1500 ————— 48s 31ms/step - accuracy: 0.8345 - loss: 0.5107 - val_accuracy: 0.9798 - val_loss: 0.0687
Epoch 2/10
1500/1500 ————— 45s 30ms/step - accuracy: 0.9670 - loss: 0.1096 - val_accuracy: 0.9858 - val_loss: 0.0466
Epoch 3/10
1500/1500 ————— 82s 30ms/step - accuracy: 0.9785 - loss: 0.0746 - val_accuracy: 0.9884 - val_loss: 0.0404
Epoch 4/10
1500/1500 ————— 82s 30ms/step - accuracy: 0.9806 - loss: 0.0641 - val_accuracy: 0.9898 - val_loss: 0.0342
Epoch 5/10
1500/1500 ————— 81s 30ms/step - accuracy: 0.9819 - loss: 0.0554 - val_accuracy: 0.9896 - val_loss: 0.0342
Epoch 6/10
1500/1500 ————— 48s 32ms/step - accuracy: 0.9838 - loss: 0.0518 - val_accuracy: 0.9908 - val_loss: 0.0301
Epoch 7/10
1500/1500 ————— 81s 31ms/step - accuracy: 0.9851 - loss: 0.0484 - val_accuracy: 0.9903 - val_loss: 0.0336
Epoch 8/10
1500/1500 ————— 80s 30ms/step - accuracy: 0.9856 - loss: 0.0452 - val_accuracy: 0.9898 - val_loss: 0.0328
Epoch 9/10
1500/1500 ————— 85s 32ms/step - accuracy: 0.9870 - loss: 0.0405 - val_accuracy: 0.9913 - val_loss: 0.0328
Epoch 10/10
1500/1500 ————— 80s 30ms/step - accuracy: 0.9887 - loss: 0.0372 - val_accuracy: 0.9930 - val_loss: 0.0281

313/313  2s 7ms/step - accuracy: 0.9902 - loss: 0.0266
Test Accuracy: 0.99

```
# Save the trained model  
model.save("handwritten_character_recognition.h5")
```

 WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is consi

